

# VIRTUAL MACHINES ARCHITECTURE

Koushik Kumar Kamala  
Computer Engineering  
San Jose State University  
San Jose, USA  
koushikkumar.kamala@sjsu.edu

**Abstract—** To analyze a virtual machines architecture from operating systems point of view. A detailed analysis about how the operating system architecture addresses the platform independence of the virtual machines and how the virtual machines behave virtual existing on one single operating system.

**Keywords—** Virtual Machine, VMware, Hypervisor, Operating System, Virtualization.

## I. INTRODUCTION

A virtual machine gives a totally verified and detached duplicate of the fundamental physical system. It receives a layered technique to achieve this goal. We need another layer over the primary presented structures to separate the physical resources and offer interface to working systems running on it. This layer is known as the Virtual Machine Monitor (VMM). The VMM is the crucial bit of the virtual machine utilization, since it plays out the understanding between the revealed hardware and virtualized essential stage: giving virtual processors, memory, and virtualized I/O devices. Since all the virtual machines share the equivalent revealed gear, the Virtual Machine Monitor should in like manner give appropriate affirmation with the objective that each virtual machine is a withdrawn propagation. The principal thought behind a virtual machine is to extract the equipment of a solitary PC (the CPU, memory, circle drives, arrange interface cards, etc) into a few distinctive execution conditions, in this way making the fantasy that each different condition is running without anyone else private PC. This idea may appear to be like the layered methodology of working framework usage (see Section 2.7.2), and here and there it is. On account of virtualization, there is a layer that makes a virtual framework on which working frameworks or applications can run.

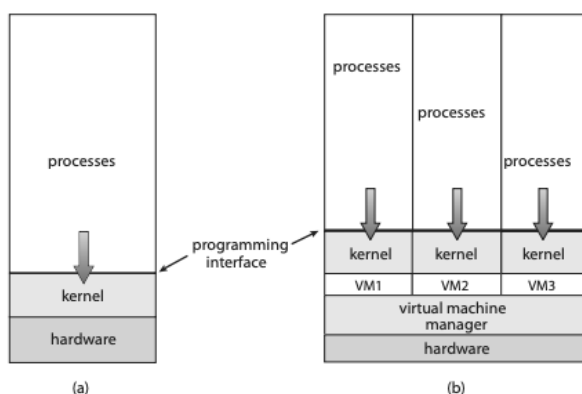


Figure (a): System models. (a) Nonvirtual machine. (b) Virtual machine.

Virtual machine usage includes a few parts. At the base is the host, the hidden equipment framework that runs the virtual machines. The virtual machine administrator (VMM) (otherwise called a hypervisor) makes and runs virtual machines by giving an interface that is indistinguishable to the host (aside from on account of paravirtualization, talked about later). Every visitor procedure is given a virtual duplicate of the host (Figure 16.1). More often than

not, the visitor procedure is in certainty a working framework. A solitary physical machine would thus be able to run numerous working frameworks simultaneously, each in its own virtual machine. Pause for a minute to take note of that with virtualization, the meaning of "working framework" indeed obscures. For instance, consider VMM programming, for example, VMware ESX. This virtualization programming is introduced on the equipment, runs when the equipment boots, and gives administrations to applications. The administrations incorporate customary ones, for example, booking and memory the executives, alongside new sorts, for example, movement of uses between frameworks. Besides, the applications are in reality visitor working frameworks. Is the VMware ESX VMM a working framework that, thus, runs other working frameworks? Unquestionably it acts like a working framework. For lucidity, nonetheless, we consider the segment that gives virtual situations a VMM. The execution of VMMs shifts extraordinarily. Alternatives incorporate the accompanying: Hardware-based arrangements that offer help for virtual machine creation and the executives through firmware. These VMMs, which are regularly found in centralized computer and vast to moderate sized servers, are commonly known as sort 0 hypervisors.

## II. HISTORY

The idea of virtual machines was first created by IBM during the 1960s to give simultaneous, intelligent access to a centralized server PC and first showed up financially on IBM centralized computers in 1972. Virtualization was given by the IBM VM working framework. This framework has advanced is as yet accessible. What's more, a significant number of its unique ideas are found in different frameworks, making it worth investigating. IBM VM370 separated a centralized computer into different virtual machines, each running its own working framework. A noteworthy trouble with the VM approach included circle frameworks. Assume that the physical machine had three plate drives however needed to help seven virtual machines. Plainly, it couldn't allot a circle drive to each virtual machine. The arrangement was to give virtual plates—named minidisks in IBM's VM working framework. The minidisks are indistinguishable to the framework's hard circles in all regards aside from size. The framework actualized each minidisk by apportioning the same number of tracks on the physical plates as the minidisk required. When the virtual machines were made, clients could run any of the working frameworks or programming bundles that were accessible on the fundamental machine.

For the IBM VM system, a user ordinarily ran CMS—a solitary user intuitive working framework. For a long time after IBM presented this innovation, virtualization stayed in its area. Most systems couldn't bolster virtualization. Be that as it may, a formal meaning of virtualization built up framework prerequisites and an objective for usefulness. The virtualization necessities expressed that: 1. A VMM gives a situation to programs that is basically indistinguishable to the first machine. 2. Projects running inside that condition show

just minor execution diminishes. 3. The VMM is in finished control of framework assets.

### III. BUILDING BLOCKS

The capacity to virtualize relies upon the highlights given by the CPU. On the off chance that the highlights are adequate, at that point it is conceivable to compose a VMM that gives a visitor domain. Something else, virtualization is unthinkable. VMMs utilize a few methods to actualize virtualization, including trap-and-emulate and paired interpretation. We talk about every one of these systems in this segment, alongside the equipment bolster expected to help virtualization. One essential idea found in most virtualization choices is the execution of a virtual CPU (VCPU). The VCPU does not execute code. Or maybe, it speaks to the condition of the CPU as the visitor machine trusts it to be. For every visitor, the VMM keeps up a VCPU speaking to that visitor's present CPU state.




Figure (b): Computer system architecture. Key implementation layers communicate vertically via the instruction set architecture (ISA), application binary interface (ABI), and application programming interface (API)

A discussion of VMs is also a discussion about PC engineering in the unadulterated sense of the term. Because VM implementations lie at architected interfaces, a noteworthy consideration in the construction of a VM is the devotion with which it implements these interfaces. Engineering, as connected to PC systems, refers to a formal specification of an interface in the system, including the intelligent conduct of resources oversight by means of the interface. Execution describes the real encapsulation of an engineering. Abstraction levels correspond to usage layers, regardless of whether in hardware or software, each associated with its own interface or engineering. Three of these interfaces at or close to the HW/SW limit—the instruction set design, the application double interface, and the application programming interface—are especially vital for VM construction

#### A. Trap-and-Emulate:

On a typical dual-mode system, the virtual machine visitor can execute just in user mode (except if additional equipment support is given). The piece, obviously, keeps running in part mode, and it isn't sheltered to permit user level code to keep

running in bit mode. Similarly, as the physical machine has two modes, notwithstanding, so should the virtual machine.

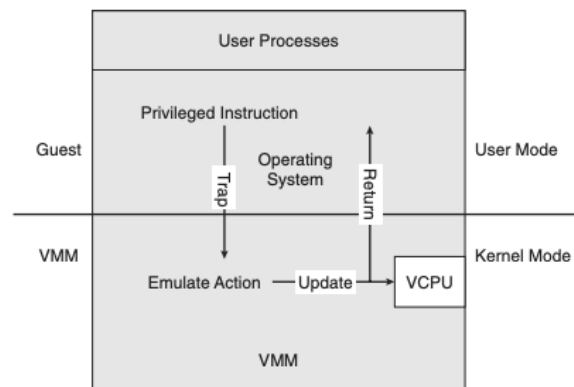


Figure (c): Trap-and-emulate virtualization implementation.

Therefore, we should have a virtual client mode and a virtual portion mode, the two of which keep running in physical client mode. Those activities that reason an exchange from client mode to piece mode on a genuine machine, should likewise make an exchange from virtual client mode virtual bit mode in the virtual machine. In what capacity can such an exchange be cultivated? The method is as per the following: When the piece in the visitor endeavors to execute a special guidance, that is a blunder (in light of the fact that the framework is in client mode) and makes a snare the VMM in the genuine machine.

The VMM picks up control and executes (or "imitates") the activity that was endeavored by the visitor bit with respect to the visitor. It at that point returns control to the virtual machine. This is known as the trap-and-emulate strategy. Most virtualization items utilize this technique to one degree or other. With special guidelines, time turns into an issue. All nonprivileged guidelines run locally on the equipment, giving indistinguishable execution to visitors from local applications. Special directions make additional overhead, notwithstanding, making the visitor run more gradually than it would locally. Likewise, the CPU is being multi programmed among numerous virtual machines, which can additionally back off the virtual machines in erratic ways.

#### B. Binary Translation Analysis

Few CPUs don't have a partition of privileged and nonprivileged directions. Tragically for virtualization implementers, the Intel x86 CPU line is one of them. The chip has kept up in reverse similarity all through its lifetime, anticipating changes that would have made virtualization less demanding through numerous ages. We should think about a case of the issue. The direction popf loads the banner register from the substance of the stack. On the off chance that the CPU is in privileged mode, the majority of the banners are supplanted from the stack. On the off chance that the CPU is in user mode, at that point just a few banners are supplanted, and others are disregarded. Since no snare is created if popf is executed in client mode, the trap-and-emulate technique is rendered pointless. Other x86 guidelines cause comparative issues.




Figure (c): Binary translation virtualization implementation

For the reasons for this exchange, we will call this arrangement of directions exceptional guidelines. The utilizing the trap-and-emulate strategy to execute virtualization on the x86 was viewed as outlandish on account of these uncommon guidelines. This beforehand inconceivable issue was illuminated with the execution of the parallel interpretation strategy. Twofold interpretation is genuinely basic in idea yet complex in execution. The fundamental advances are as per the following: 1. On the off chance that the visitor VCPU is in client mode, the visitor can run its guidelines locally on a physical CPU. 2. On the off chance that the visitor VCPU is in kernel mode, at that point the visitor trusts that it is running in kernel mode. The VMM analyzes each guidance the visitor executes in virtual piece mode by perusing the following couple of guidelines that the visitor will execute, in light of the visitor's program counter. Directions other than extraordinary guidelines are run locally.

#### IV. HARDWARE SUPPORT

Without hardware support, virtualization would be unthinkable. The more equipment is accessible inside a framework, the more component rich and stable the virtual machines can be and the better they can perform. In the Intel x86 CPU family, Intel included new virtualization support in progressive ages (the VT-x guidelines) starting in 2005. Presently, parallel interpretation is never again required. Truth be told, all real broadly useful CPUs are giving broadened measures of equipment support for virtualization. For instance, AMD virtualization innovation (AMD-V) has showed up in a few AMD processors beginning in 2006. It characterizes two new methods of activity—host and guest—in this way moving from a double mode to a multimode processor. The VMM can empower host mode, characterize the attributes of every guest virtual machine, and after that change the system to guest mode, passing control of the system to a guest working framework that is running in the virtual machine.

In guest mode, the virtualized working system supposes it is running on local equipment and sees whatever gadgets are incorporated into the host's meaning of the guest. In the event that the guest endeavors to get to a virtualized asset, at that point control is passed to the VMM to deal with that communication. The usefulness in Intel VT-x is comparable, giving root and non-root modes, proportionate to host and guest modes. Fundamentally, these CPUs execute

settled page tables in equipment to enable the VMM to completely control paging while the CPUs quicken the interpretation from virtual to physical locations. The NPTs include another layer, one speaking to the guest perspective on consistent to-physical location interpretation. The CPU page-table strolling capacity incorporates this new layer as important, strolling through the visitor table to the VMM table to locate the physical location wanted.

#### V. LIFE CYCLE

These parameters generally incorporate the quantity of CPUs, measure of memory, organizing subtleties, and capacity subtleties that the VMM will consider while making the visitor. The VMM at that point makes the virtual machine with those parameters. On account of a sort 0 hypervisor, the assets are generally devoted. In this circumstance, if there are not two virtual CPUs accessible and unallocated, the creation ask for in our precedent will fall flat. For other hypervisor types, the assets are committed or virtualized, contingent upon the sort. Positively, an IP address can't be shared, yet the virtual CPUs are normally multiplexed on the physical CPUs. Essentially, memory the board for the most part includes allotting more memory to visitors than really exists in physical memory. At long last, when the virtual machine is never again required, it tends to be erased.

At the point when this occurs, the VMM first opens up any utilized plate space and after that evacuates the arrangement related with the virtual machine, basically overlooking the virtual machine. These means are very straightforward contrasted and building, designing, running, and evacuating physical machines. Making a virtual machine from a current one can be as simple as tapping the "clone" button and giving another name and IP address. This simplicity of creation can prompt virtual machine spread, which happens when there are such a large number of virtual machines on a framework that their utilization, history, and state end up confounding and hard to follow.




Figure (d): Multiplexing




Figure (e): Suspension




Figure (f): Provision



Figure (g): Life Migration

These VM operations empower a virtual machine to be provisioned to any accessible hardware stage. They make it adaptable to port distributed application executions. Moreover, the VM approach will significantly upgrade the use of server resources. Various server functions can be consolidated on the same hardware stage to accomplish higher system proficiency. This will dispense with server sprawl by means of sending of systems as VMs. These VMs move transparency to the shared hardware. As indicated by a case by VMWare, the server use could be increased from current 5-15% to 60-80%.

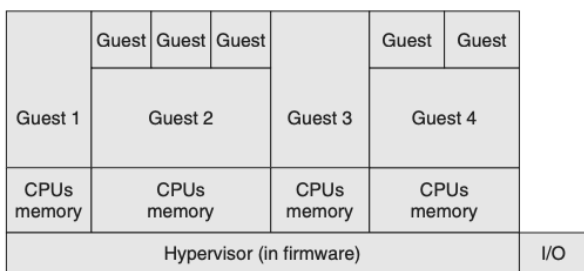


Figure (h): Type 0 Hypervisor

Type 0 hypervisors have existed for a long time as "partitions" and "domains" etc. They are an equipment highlight, and that brings its own positives and negatives. Operating Systems need do nothing unique to exploit their highlights. The VMM itself is encoded in the firmware and stacked at boot time. Thus, it stacks the visitor pictures to keep running in each segment. The list of capabilities of a type 0 hypervisor will in general be littler than those of alternate kinds since it is executed in equipment. For instance, a framework may be part into four virtual frameworks, each with committed CPUs, memory, and I/O gadgets. Every guest trusts that it has committed equipment since it does, streamlining numerous usage subtleties. Since type 0

virtualization is extremely near crude equipment execution, it ought to be considered independently from alternate techniques talked about here. A type 0 hypervisor can run numerous guest OS (one in every equipment segment). Those visitors, since they are running on basic hardware, can thus be VMMs. Basically, the visitor working frameworks in a type 0 hypervisor are local working frameworks with a subset of equipment made accessible to them. Thus, each can have its very own visitor working frameworks. Different sorts of hypervisors for the most part can't give this virtualization-inside virtualization usefulness.




Figure (i): Type 1 Hypervisor

Type 1 hypervisors are generally found in organization server and data. They are extraordinary reason operating systems that run locally on the hardware, but instead notwithstanding running on standard hardware, they can keep running on type 0 hypervisors, yet not on other type 1 hypervisors. Whatever the stage, guests for the most part don't realize they are running on anything besides the local hardware. By utilizing type 1 hypervisors, an imperative advantage is the capacity to solidify all the more operating systems and applications onto less systems. For instance, instead of having ten systems running at 10 percent usage each, a server farm may have one server deal with the whole burden. In the event that use expands, guests and their applications can be moved to less-stacked systems live, without interference of administration. Utilizing previews and cloning, the framework can spare the conditions of guests and copy those states, an a lot less demanding assignment than reestablishing from reinforcements or introducing physically or through contents and instruments. The cost of this expanded sensibility is the expense of the VMM and the need to adapt new administration devices and techniques.

This design adds an extra virtualization layer, consisting of two sub layers (VMs and type-1 hypervisor), to an engineering that was at that point made complex by having numerous cores. By giving simulated hardware environments, the VMs empower the use of various operating systems. Applications share VMs, the hypervisor, and cores as shared resources, consequently giving single points of disappointment and possible interference paths. Interference can happen when the execution of one application running in one VM affects the execution of another application running in a second VM by damaging




Figure (j): Type 2 Hypervisor (Pure and Practical)

Type 2 hypervisors are less intriguing to us as operating-system adventurers, on the grounds that there is almost no operating-system association in these application level VMMs, is just another procedure run and overseen by the host, and even the host does not realize virtualization is occurring inside the VMM. Type 2 hypervisors have limits not related with a portion of alternate types. As is frequently the situation, the constraints of type 2 hypervisors additionally give a few advantages. They keep running on an assortment of universally useful operating systems and running them requires no progressions to the host operating framework. An understudy can utilize a type 2 hypervisor, for instance, to test a non-local operating framework without supplanting the local operating framework.

## VI. COMPONENTS

### a. CPU Scheduling

A system with virtualization, even a single-CPU system, oftentimes acts like a multiprocessor system. The virtualization software presents at least one virtual CPUs to every one of the virtual machines running on the system and afterward schedules the use of the physical CPUs among the virtual machines. The significant differences among virtualization technologies make it hard to summarize the impact of virtualization on scheduling. Coming to VMM scheduling, it has various physical CPUs accessible and various threads to keep running on those CPUs. The threads can be VMM threads or guest threads. Guests are designed with a specific number of virtual CPUs at creation time, and that number can be adjusted for the duration of the life of the VM. At the point when there are sufficient CPUs to distribute the requested number to every guest, VMM can regard the CPUs as devoted and schedule just a given guest's threads on that guest's CPUs

### b. Memory Management

Productive memory usage in multi-purpose operating systems is important for performance. In virtualized environments, there are more users of memory prompting more pressure on memory use. Adding to this, VMMs regularly overcommit memory, so that the absolute memory with which guests are arranged exceeds the measure of memory that physically exists in the system. The additional requirement for productive memory use is not lost on the implementers of VMMs, who take incredible measures to ensure the ideal use of memory.

### c. I/O

Hypervisors account for some part in I/O and can be less worried about precisely representing the hidden hardware to their guests. Because of all the variety in I/O devices, operating systems are used to managing differing and adaptable I/O mechanisms. Operating systems have a device-driver mechanism that provides a uniform interface to the operating system whatever the I/O device. Device-driver interfaces are designed to permit outsider hardware manufacturers to give device drivers associating their devices to the operating system. Usually, device drivers can be powerfully stacked and emptied. Virtualization takes favorable position of such inherent adaptability by giving specific virtualized devices to guest operating systems.

Interference occurs when software executing on one VM impacts the conduct of software executing on different VMs. This interference includes failures of both spatial isolation (because of shared memory access) and disappointment of worldly isolation (because of interference delays as well as penalties). As is the case with multicore processors, the quantity of interference paths increases very quickly with the quantity of VMs. Consequently, an exhaustive analysis of all interference paths is frequently impossible. The difficulty of exhaustive analysis necessitates the selection of representative interference paths while dissecting isolation. The accompanying outline uses the shading red to show three possible interference paths including VMs and their hypervisor between pairs of applications including six shared resources.

### d. Storage Management

Virtualized environments approach the storage management differently in contrast to local operating systems. Indeed, even the standard multiboot strategy for slicing the root disk into partitions, installing a boot administrator in one partition, and installing each other operating system in another partition is not sufficient, because partitioning has limits that would keep it from working for tens or hundreds of virtual machines. By and by, the solution to this issue depends on the type of hypervisor. Type 0 hypervisors do will in general permit root disk partitioning, mostly because these systems will in general run less guests than different systems. On the other hand, they may have a disk chief as a component of the control partition, and that disk director provides disk space to alternate partitions.

### e. Live Migration

One feature not found in general-purpose operating systems but found in type 0 and type 1 hypervisors is the live migration of a running guest from one system to another. We mentioned this capability earlier. Here, we explore the details of how live migration works and why VMMs have a relatively easy time implementing it while general-purpose operating systems, in spite of some research attempts, do not. First, consider how live migration works. A running guest on one system is copied to another system running the same VMM. The copy occurs with so little interruption of service that users logged in to the guest, and network connections to the guest, continue without noticeable impact. This rather astonishing ability is very powerful in resource management



and hardware administration. After all, compare it with the steps necessary without virtualization: warning users, shutting down the processes, possibly moving the binaries, and restarting the processes on the new system, with users only then able to use the services again. With live migration, an overloaded system can have its load decreased live with no discernible disruption. Similarly, a system needing hardware or system changes (for example, a firmware upgrade, hardware addition or removal, or hardware repair) can have guests migrated off, the work done, and guests migrated back without noticeable impact on users or remote connections.

One component not found in usual OS is the live migration of a running guest starting with one system then onto the next. A running guest on one system is duplicated to another system running the same VMM. The duplicate occurs with so little intrusion of service that users signed in to the guest, and system connections to the guest, proceed without recognizable effect. This ability is incredible in resource the executives and hardware administration. All things considered, contrast it and the steps necessary without virtualization: cautioning users, shutting down the processes, possibly moving the binaries, and restarting the processes on the new system, with users at exactly that point ready to use the services once more. With live migration, an over-burden system can have its heap decreased live with no discernible disruption. Similarly, a system requiring hardware or system changes can have guests moved off, the work done, and guests moved back without recognizable effect on users or remote connections.

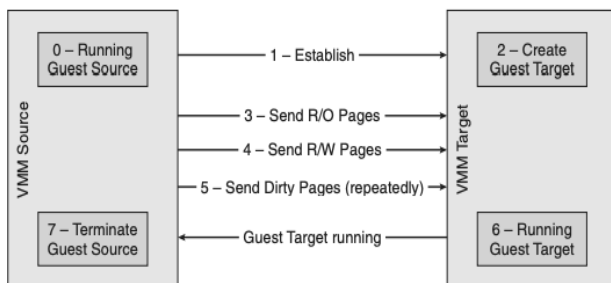


Figure (k): Live migration of a guest between two servers.

A restriction of live migration is that no disk state is transferred. One reason live migration is possible is that most of the guest's state is kept up inside the guest—for instance, open record tables, system-call state, part state, etc. Because disk I/O is so much slower than memory access and used disk space is usually a lot bigger than used memory, disks associated with the guest can't be moved as a component of a live migration. Or maybe, the disk must be remote to the guest, accessed over the system. In that case, disk access state is kept up inside the guest, and system connections are the only things that are in any way important to the VMM. The system connections are kept up amid the migration, so that remote disk access is not disturbed.

## VII. SUMMARY

Virtualization is a strategy for giving a guest a copy of a system's hardware. Different guests can keep running on a given system, each trusting it is the local operating system in full control of the system. Virtualization started as a strategy to permit IBM to segregate users and furnish them with their own execution environments on IBM mainframes. Since at that point, with improvements in system and CPU execution and through inventive software techniques, virtualization has turned into a typical element in server farms and even on personal computers. Because of the ubiquity of virtualization, CPU designers have added features to support virtualization. This snowball impact is probably going to proceed, with virtualization and its hardware support increasing after some time. Type 0 virtualization is executed in the hardware and requires modifications to the operating system to ensure appropriate activity. These modifications offer a case of paravirtualization, in which the operating system is not oblivious in regard to virtualization but rather instead has features added and algorithms changed to improve virtualization's features and execution. In Type 1 virtualization, a host virtual machine screen (VMM) provides the earth and features expected to make, run, and destroy guest virtual machines. Every guest includes the majority of the software commonly associated with a full local system, including the operating system, device drivers, applications, user accounts, etc. Type 2 hypervisors are simply applications that kept running on other operating systems, which don't have the foggiest idea about that virtualization is occurring. These hypervisors despise hardware or host support so must play out all virtualization activities with regards to a process. Different facilities that are similar to virtualization yet don't meet the full meaning of recreating hardware precisely are also normal. Programming condition virtualization is a piece of the design of a programming language. The language specifies a containing application in which programs run, and this application provides services to the programs. Copying is used when a host system has one design and the guest was accumulated for an alternate engineering. Each instruction the guest wants to execute must be translated from its instruction set to that of the local hardware. In spite of the fact that this strategy involves some perform punishment, it is adjusted by the usefulness of having the capacity to run old programs on more up to date, inconsistent hardware or run games designed for old consoles on current hardware. Executing virtualization is testing, especially when hardware support is negligible.

## VIII. REFERENCES

- [1] Operating Systems, Abraham Silberschatz, Peter Baer Galvin, Greg Gagne : Virtual Machines
- [2] Donald Firesmith : [https://insights.sei.cmu.edu/sei\\_blog/2017/09/virtualization-via-virtual-machines.html](https://insights.sei.cmu.edu/sei_blog/2017/09/virtualization-via-virtual-machines.html)
- [3] Architecture Of Virtual Machines\* R.P. Goldberg: <http://www.cse.psu.edu/~buul/teaching/spring06/papers/goldberg.pdf>
- [4] Class Slides